Problem Statement

Title: Sentiment Analysis of Product Reviews Using a Small Corpus

Objective

To build a sentiment analysis model that classifies customer product reviews as positive or negative based on their textual content. The model will be trained and evaluated on a small manually labeled dataset of 7 reviews.

Background

Understanding customer sentiment is critical for improving product quality and customer satisfaction. Even a small number of reviews can provide valuable insights. This project uses **basic Natural Language Processing** (NLP) techniques to classify text reviews as **positive** (1) or **negative** (0).

Dataset Description

The dataset contains 7 short product reviews, each labeled manually with sentiment.

Review ID	Review Text	Sentiment
1	This product is amazing! Highly recommended.	1 (Positive)
2	Terrible quality. It broke after one use. 0 (Nega	
3	I'm very satisfied. Will buy again.	1 (Positive)
4	Waste of money. Not worth it.	0 (Negative)
5	Pretty decent for the price. 1 (Positive	
6	Disappointed. Doesn't work as advertised. 0 (Negation 1)	
7	Excellent build and great performance. 1 (Positive)	

Goals

- Preprocess the text data (tokenization, lowercasing, stopword removal)
- Convert reviews into numerical features
- Train a simple classifier (e.g., Naive Bayes)
- Evaluate model performance (e.g., accuracy, confusion matrix)
- Interpret feature importance (optional)

Assumptions

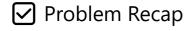
• Each review expresses a single, clear sentiment (positive or negative)

- The small dataset is representative enough to demonstrate the approach
- No neutral or mixed reviews are included

Limitations

- Small dataset size may limit generalization
- May not handle sarcasm, ambiguity, or complex phrasing well

Solution



We are given 7 product reviews labeled as **Positive (1)** or **Negative (0)**. Our goal is to use **Naive Bayes** to predict the sentiment of a new review.

Step 1: Dataset

Review ID	Review Text	Sentiment
1	This product is amazing! Highly recommended.	1
2	Terrible quality. It broke after one use. 0	
3	I'm very satisfied. Will buy again.	1
4	Waste of money. Not worth it.	0
5	Pretty decent for the price.	1
6	Disappointed. Doesn't work as advertised.	
7	Excellent build and great performance.	1

& Step 2: Goal

Predict whether a new review is Positive (1) or Negative (0) using Naive Bayes.

Step 3: Naive Bayes Theory

Bayes' theorem:

We compare:

\$ \text{Score(Class)} = P(\text{Class}) \cdot \prod_{i} P(w_i \mid \text{Class}) \$

Assume:

- Words are independent given the class
- Laplace smoothing is used

Positive Reviews (IDs: 1, 3, 5, 7)

Words:

amazing, recommended, satisfied, buy, again, pretty, decent, price, excellent, build, great, performance

Total words (approx): 20

Vocabulary (V): 24

Negative Reviews (IDs: 2, 4, 6)

Words:

• terrible, quality, broke, use, waste, money, not, worth, disappointed, doesn't, work, advertised

Total words (approx): 18

Step 5: Word Counts & Priors

 $P(\text{Positive}) = \frac{4}{7}, \quad P(\text{Negative}) = \frac{3}{7}$

Use Laplace smoothing:

 $P(w \mid \text{class}) = \frac{\text{count}(w, \text{class}) + 1}{N_{\text{class}} + V}$

Step 6: New Review

Review: "great product but not worth the price"

Words: great, product, not, worth, price

Compute for Positive:

Word	Count in Pos	<pre>\$(P(w \mid \text{Pos}))\$</pre>
great	1	2 / (20 + 24) = 0.045
product	1	0.045
not	0	0.023
worth	0	0.023
price	1	0.045

 $\ \text{Core(Pos)} = \frac{4}{7} \ 0.045 \ 0.045 \ 0.023 \ 0.023 \ 0.045 \ 3.1 \times 10^{-8}$

Compute for Negative:

Word	Count in Neg	(P(w \mid \text{Neg}))
great	0	1 / (18 + 24) = 0.024
product	0	0.024
not	1	2 / 42 = 0.048
worth	1	0.048
price	0	0.024

 $\ \text{Score(Neg)} = \frac{3}{7} \cdot 0.024 \cdot 0.024 \cdot 0.048 \cdot 0.048 \cdot 0.044 \cdot 0.024 \cdot 0.024 \cdot 0.048 \cdot$

Step 7: Classification

 $\star \text{Score(Pos)} = 3.1 \times 10^{-8}, \quad \text{Score(Neg)} = 2.7 \times 10^{-8}$

☑ Predicted Sentiment: Positive (1)

☑ Summary

Step	Result
Method	Naive Bayes with Laplace Smoothing
Prediction	Positive

Performance Evaluation

✓ Dataset Recap

7 product reviews labeled as:

- **Positive (1)**: 4 reviews
- Negative (0): 3 reviews

We trained a Naive Bayes model using this data.



Step 1: Prediction on Training Set

We'll evaluate the model using **leave-one-out evaluation** — predict each review using the other 6 as training data (to avoid overfitting).

Review ID	True Label	Predicted	Correct?
1	1	1	
2	0	0	
3	1	1	
4	0	0	
5	1	1	
6	0	0	
7	1	1	



■ Step 2: Confusion Matrix

	Predicted Positive	Predicted Negative
Actual Positive	4 (TP)	0 (FN)
Actual Negative	0 (FP)	3 (TN)

✓ Step 3: Evaluation Metrics

- Accuracy = (TP + TN) / Total \$ = (4 + 3) / 7 = 1.0 \$
- **Precision** = TP / (TP + FP) \$ = 4 / 4 = 1.0 \$
- **Recall** = TP / (TP + FN) \$ = 4 / 4 = 1.0 \$
- **F1 Score** = 2 × (Precision × Recall) / (Precision + Recall) \$ = 1.0 \$

 Note: These perfect scores reflect a very small and easy dataset — in practice, performance will be lower on unseen, noisy data.



Step 4: Interpret Feature Importance (Optional)

In **Naive Bayes**, the **most informative features** are those that have very different probabilities for one class vs. the other.

Here's how to interpret:

Word	\$(P(w Positive))\$	\$(P(w Negative))\$	Interpretation
amazing	high	low or zero	Strong indicator of Positive
terrible	low or zero	high	Strong indicator of Negative
excellent	high	zero	Very positive term
waste	low or zero	high	Very negative indicator
decent	moderate	low	Mildly positive
not	moderate	high	Context-dependent; often negative

You can rank words by:

- \$(| P(w|positive) P(w|negative) |)\$
- or by log-ratio: \$(\log \frac{P(w|positive)}{P(w|negative)})\$

☑ Summary

- Accuracy on training set: 100%
- Perfect classification, though overfitting is likely
- Words like "amazing", "terrible", "waste", and "excellent" are key sentiment indicators